

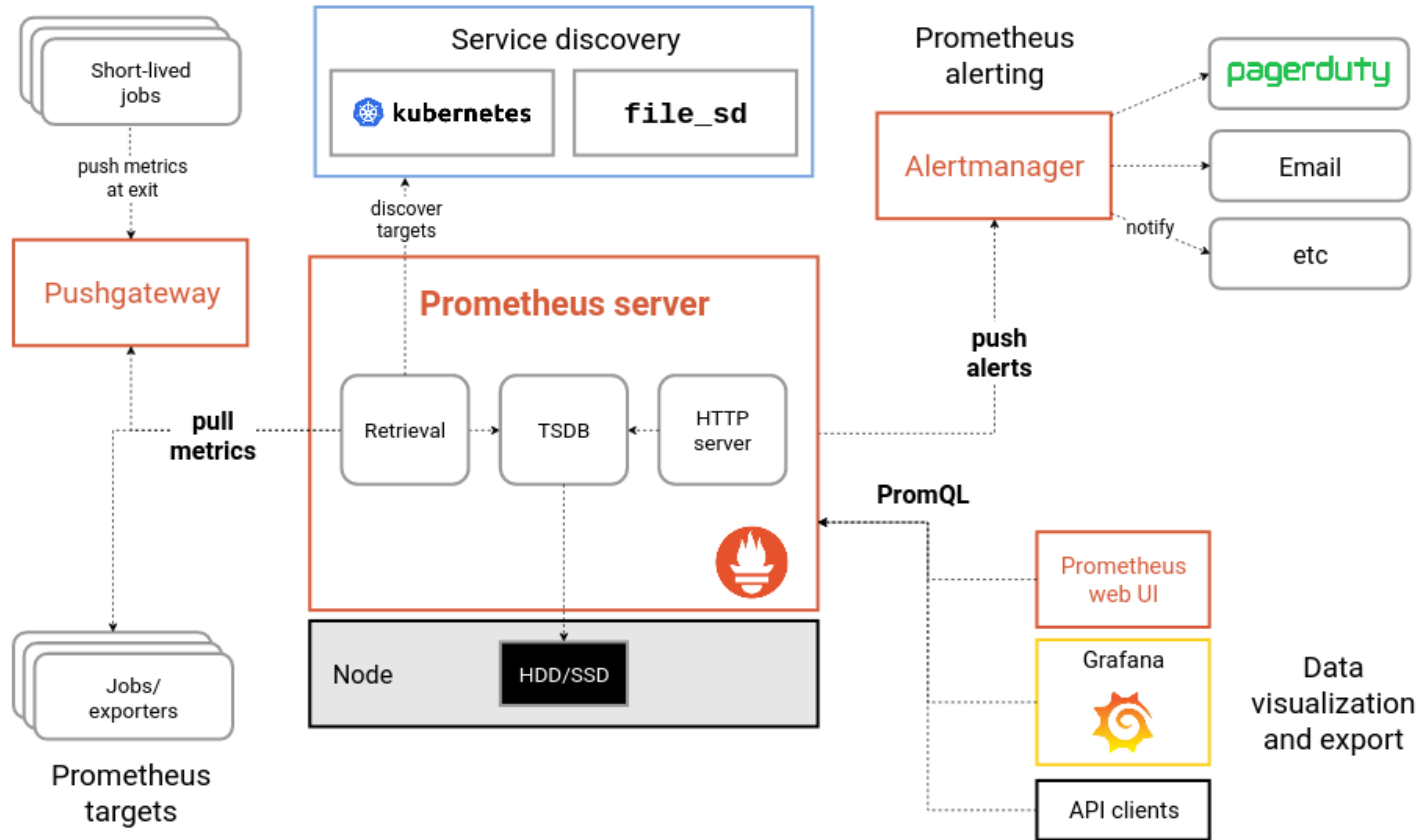
Self-healing with prometheus alerts

Dario Maiocchi <dmaiocchi@suse.com>



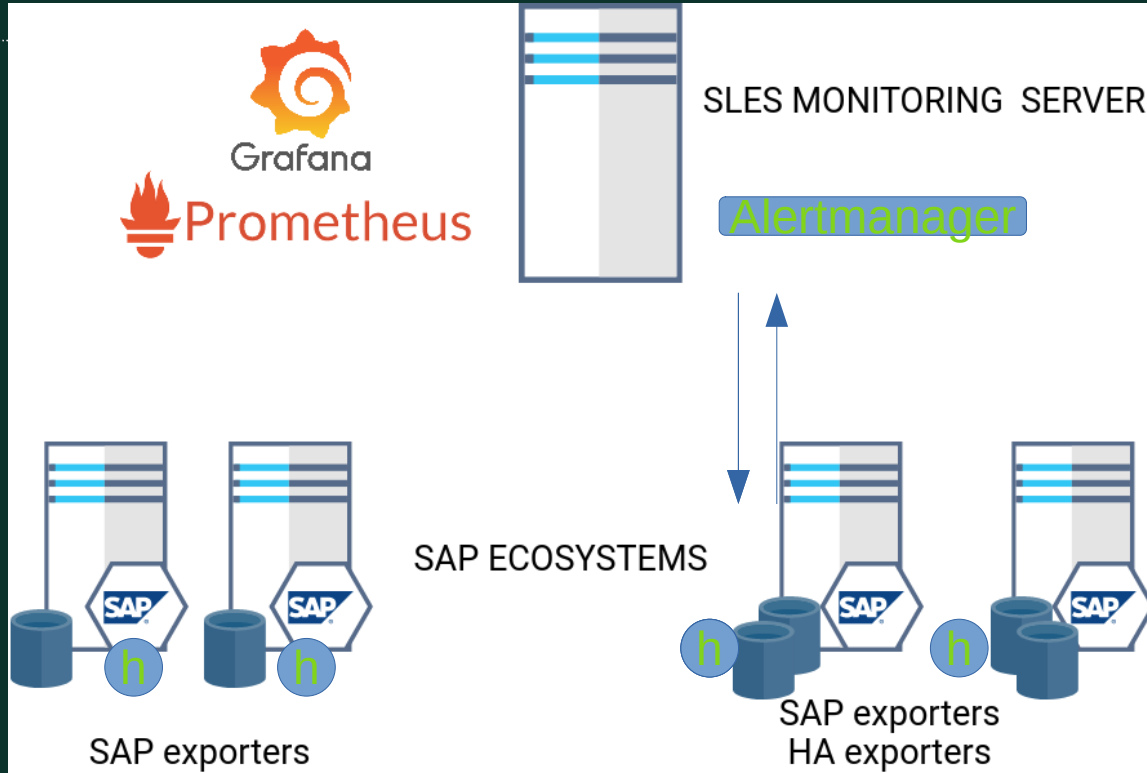
Warning:

the implementation here is for POC purposes.



USE

This is a possible architecture. (simplest)



h

H= handler. Each server has a new binary, the handler which will handle the alerts.

For technical details, infos check my confluence page

The image is a composite of two screenshots. The left screenshot shows a Confluence user profile for 'Dario Maiocchi' with a sidebar menu containing 'Profile', 'Pages', 'Blog', 'SPACE SHORTCUTS', and 'PAGE TREE'. The right screenshot shows a Confluence page titled 'Self-healing HA SAP clusters:' with a list of TODO items and a section for 'DOMAIN/HIGH-LEVEL PROBLEMS'. Below this, there is a GitHub repository view for 'MalloZup / hasap-alerts-handler' showing the 'Code' tab with 7 issues, pull requests, and actions.

Dashboard / Dario Maiocchi's Home

Self-healing HA SAP clusters:

Created by Dario Maiocchi, last modified on 2020-09-29

- TODO:
- DOMAIN/HIGH-LEVEL PROBLEMS/features:
 - Monitor and self-heal the HANA filesystem:
 - High-level description:
 - Technical implementation:
 - Internal project at suse which related:
- TECHNICAL PROBLEMS:
 - Component:
 - Design:

TODO:

- clarify why <https://docs.microsoft.com/en-us/azure/virtual-machines/workloads/sap/hana-vm-troubleshoot-scale-out-ha-on-sles#failover-or-takeover> state this complex operations.
- Clarify if XML pacemaker is an array or an element for resource

DOMAIN/HIGH-LEVEL PROBLEMS

Monitor and self-heal the HANA filesystem:

MalloZup / hasap-alerts-handler

<> Code ! Issues 7 🔗 Pull requests ⏮ Actions

🔗 master ▾ 🔗 1 branch 🔗 2

RFC will come later. (didn't had time)

Features:

- Secure (https) possible.
- Self-heal on specific alerts, based on declarative API
- Raise alerts back to alertmanager in case selfheal fail
(can be used for other purposes, like monitor config file etc)

Declarative api

```
- name: hana filesystem monitoring
  rules:
    # all critical severity will be handled as self-healing
    - alert: HanaFileSystemFull
      expr: ((node_filesystem_size_bytes{mountpoint="/hana"} - node_filesystem_size_bytes{mountpoint="/hana"})) / node_filesystem_size_bytes{mountpoint="/hana"} > 0.95
      labels:
        severity: critical
        selfhealing: true
        component: hana
      annotations:
        summary: Hana file systems is full more then %95 percent
```

Declarative api

```
- name: hana filesystem monitoring
rules:
# all critical severity will be handled as self-healing
- alert: HanaFileSystemFull
  expr: ((node_filesystem_size_bytes{mountpoint="/hana"} - node_filesystem_size_bytes{mountpoint="/hana"})) / node_filesystem_size_bytes{mountpoint="/hana"} > 0.95
  labels:
    severity: critical
    selfhealing: true
    component: hana
  annotations:
    summary: Hana file systems is full more then %95 percent
```

labels play a role on the handler

this will be documented clearly

severity:

define how severe is an alert.

```
- name: hana filesystem monitoring
  rules:
    # all critical severity will be handled as self-healing
    - alert: HanaFileSystemFull
      expr: ((node_filesystem_size_bytes{mountpoint="/hana"} - node_filesystem_size_bytes{mountpoint="/hana"})) / node_filesystem_size_bytes{mountpoint="/hana"} > 0.95
      labels:
        severity: critical
        selfhealing: true
        component: hana
      annotations:
        summary: Hana file systems is full more then %95 percent
```

selfhealing:

```
- name: hana filesystem monitoring
  rules:
    # all critical severity will be handled as self-healing
    - alert: HanaFileSystemFull
      expr: ((node_filesystem_size_bytes{mountpoint="/hana"} - node_filesystem_size_bytes{mountpoint="/hana"})) / node_filesystem_size_bytes{mountpoint="/hana"} > 0.95
      labels:
        severity: critical
        selfhealing: true
        component: hana
      annotations:
        summary: Hana file systems is full more then %95 percent
```

Disable selfhealing or enable it.

If omitted is disabled on that alert.



```
- name: hana filesystem monitoring
  rules:
    # all critical severity will be handled as self-healing
    - alert: HanaFileSystemFull
      expr: ((node_filesystem_size_bytes{mountpoint="/hana"} - node_filesystem_size_bytes{mountpoint="/hana"})) / node_filesystem_size_bytes{mountpoint="/hana"} > 0.95
      labels:
        severity: critical
        selfhealing: true
        component: hana
      annotations:
        summary: Hana file systems is full more then %95 percent
```

component:

Influence routing of alerts

Alertermanager

Route and filter alert,
send them to various
handlers

```
route:
  receiver: 'default-receiver'
  group_wait: 10s
  group_interval: 10s
  repeat_interval: 10s
  group_by: [alertname]
  # All alerts that do not match the following child routes
  # will remain at the root node and be dispatched to 'default-receiver'.

  routes:
    # All alerts with component netweaver or hana are dispatched to sap-hook.
    - receiver: 'hana-hook-receiver'
      # overwrite default root value
      # group_wait: 10s
      match_re:
        component: hana
      # All alerts with the component label match this sub-route.
    - receiver: 'ha-hook-receiver'
      # group_by: [product, environment]
      match:
        component: sbd
```

~



SUSE

Templates are possible

```
global:
  slack_api_url: '<slack_webhook_url>'

route:
- receiver: 'slack-notifications'
  group_by: [alertname, datacenter, app]

receivers:
- name: 'slack-notifications'
  slack_configs:
    - channel: '#alerts'
      text: '{{ template "slack.myorg.text" . }}'

templates:
- '/etc/alertmanager/templates/myorg.tmpl'
```

<https://prometheus.io/blog/2016/03/03/custom-alertmanager-templates/>

```

global:
receivers:
  - name: default-receiver
    webhook_configs:
      - url: "http://10.162.29.223:9999/hooks-default"
      - url: "http://10.162.32.38:9999/hooks-default"
      # this is for debugging
      - url: "http://10.163.16.126:9999/hooks-default"

  - name: ha-hook-receiver
    webhook_configs:
      - url: "http://10.162.29.223:9999/hooks-ha"
      # this is for debugging
      - url: "http://10.163.16.126:9999/hooks-default"

  - name: hana-hook-receiver
    webhook_configs:
      - url: "http://10.162.29.223:9999/hooks-sap"
      # this is for debugging
      - url: "http://10.163.16.126:9999/hooks-sap"

route:
  receiver: 'default-receiver'
  group_wait: 10s
  group_interval: 10s
  repeat_interval: 10s
  group_by: [alertname]
  # All alerts that do not match the following child routes
  # will remain at the root node and be dispatched to 'default-receiver'.

  routes:
    # All alerts with component netweaver or hana are dispatched to sap-hook.
    - receiver: 'hana-hook-receiver'
      # overwrite default root value
      # group_wait: 10s
      match_re:
        component: hana
    # All alerts with the component label match this sub-route.
    - receiver: 'ha-hook-receiver'
      # group_by: [product, environment]
      match:
        component: sbd

```



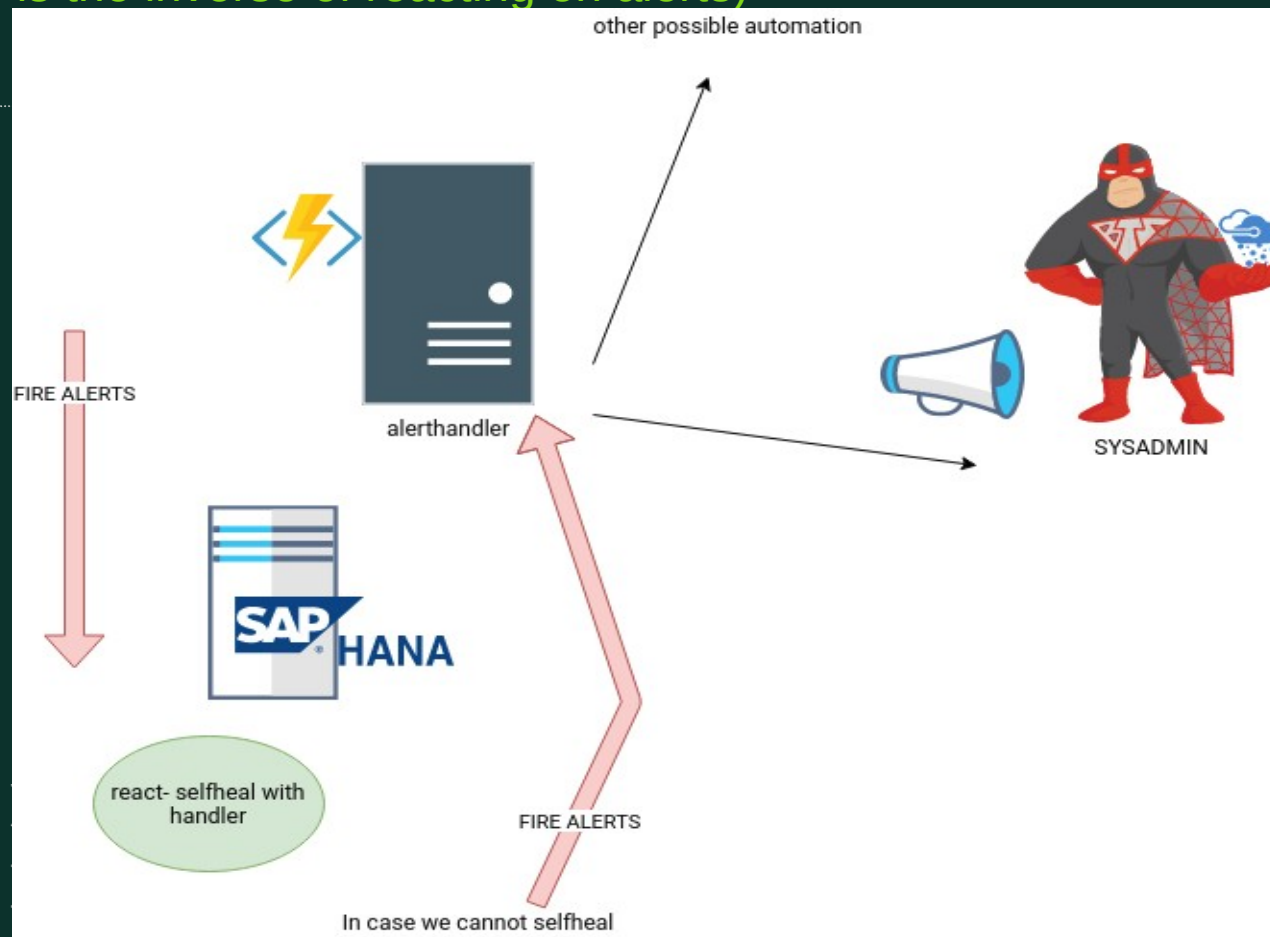
From an application codebase we can fire alerts.
This is needed for the handlers when they cannot selfhealf

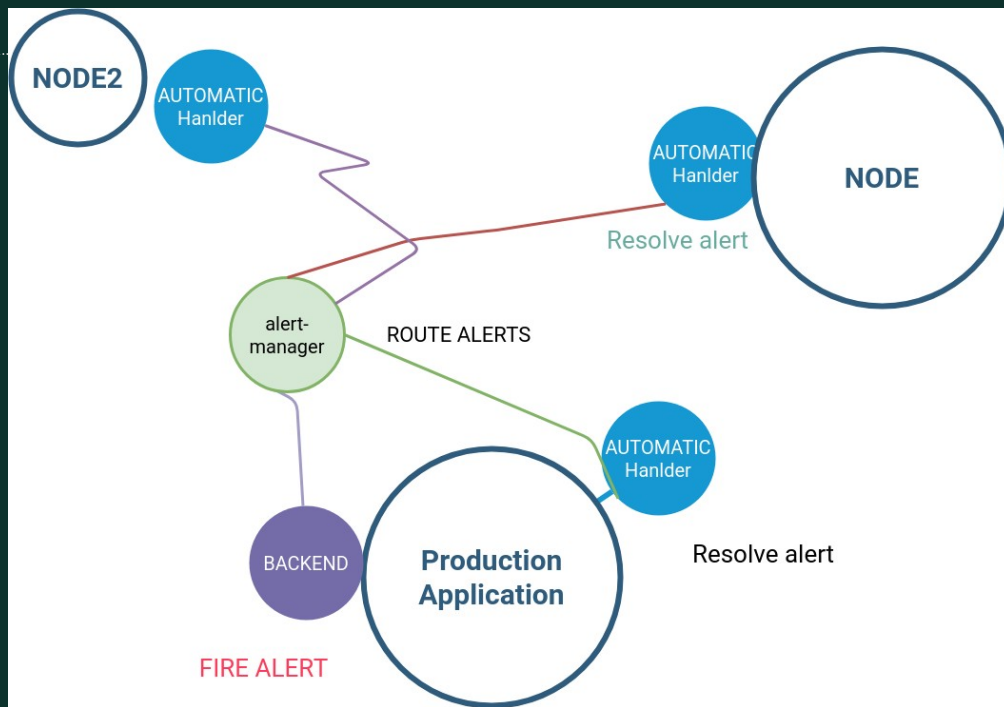
but it can be reused for other purposes, like monitoring configuration files change etc..

1) Fire alert from your application

```
if diskFull == True { // your critical condition
    log.Error("full disk ... etc")
    // ..
    var a *AlertFire
    a = new(AlertFire)
    a.Status = "firing"
    a.Labels.Alertname = "FOO-ALERT"
    a.Labels.Component = "unit-test component"
    a.Labels.Severity = "critical"
    a.Labels.Instance = "test instance"
    a.Annotations.Summary = "just a test"
    a.GeneratorURL = "unit-test"
    // this is your prometheus alertmanager server
    a.sendAlert("http://10.162.31.2:9093/api/v1/alerts")
}
```


Firing alerts from application codebase. (This is the inverse of reacting on alerts)





For details:

<https://mallozup.github.io/posts/selfhealing-systems-part2/>

DEMO



Thank you.

